

Mitra AI Product Requirements Document

Overview

Mitra AI is a mobile first UPSC preparation platform designed as a success operating system for aspirants who struggle with content overload, weak feedback loops, motivation decay, and limited peer accountability. The product direction is based on the case study brief and on the actual prototype evolution provided here: the original MVP concept, the AI generated design prompt, and the final Figma generated app experience with onboarding, dashboard, mentor, notes, progress, current affairs, quiz, pods, and settings.

This document consolidates those artifacts into a fuller product manager style pre prototype and prototype readiness document. It focuses only on capabilities that were actually built or explicitly defined in the supplied materials, while expanding the reasoning, flows, and decision trade offs in a format that is closer to a detailed product document than a short pitch note.

Product Vision

The case study identifies a core market gap: the exam preparation ecosystem has abundant content, but insufficient direction, feedback, and behavioral support. Mitra AI addresses that gap by combining an AI mentor, guided study flows, progress visibility, relevant current affairs, and lightweight accountability mechanisms inside a single mobile product rather than spreading the journey across disconnected PDFs, videos, chat groups, and test apps.

The vision is not to out compete existing players on raw content volume. The goal is to help a UPSC aspirant decide what to study, measure whether progress is real, stay motivated over a long preparation cycle, and convert scattered effort into daily execution. That focus aligns with the case study requirement to build a success operating system instead of another generic content library.

Problem Framing

The product starts from a behavioral problem, not a content problem. UPSC aspirants already have access to a large amount of study material through YouTube, Telegram channels, downloaded PDFs, coaching portals, and note sharing communities, yet outcomes remain poor and dropout rates remain high.

Several recurring pain points are visible in the case study and are echoed in the prototype concept:

- Content abundance but direction scarcity, where users have too many resources but no clear syllabus aware sequence for what to do next
- The collection trap, where time is spent gathering notes and PDFs rather than studying them

- The feedback gap, where users cannot easily tell whether they are on track today because meaningful feedback comes too late
- Motivation decay and burnout during long, isolating self study cycles
- Current affairs overload, where information is available but not mapped to exam relevance
- Connectivity constraints for learners in lower bandwidth environments, which makes heavy content experiences less accessible
- Limited peer interaction, reducing accountability and healthy competition in a long preparation journey
- Static question banks that may not feel aligned with changing exam patterns

These problems create an important product insight. The real friction is not just lack of information. It is poor decision support, weak daily feedback, fragmented learning surfaces, and inconsistency in study behavior. That is why the MVP centers on a mentor and operating system model rather than starting with a giant lecture repository.

Target User Identification

The initial target user should be narrow enough for a seed stage team to serve well. The chosen primary segment is UPSC Civil Services aspirants, especially users who are self directed, digitally active, and trying to build structure without fully relying on a physical coaching center.

The most relevant first users from the provided prototype are:

- UPSC aspirants preparing seriously for the Civil Services exam
- First attempt and second attempt candidates who still need structure, confidence, and feedback
- Tier 2 and Tier 3 city learners who may have uneven access to coaching and internet quality
- Working professionals or part time learners with limited daily study hours
- Self study aspirants who rely heavily on online resources and need curated guidance instead of more content
- Community oriented learners who benefit from streaks, peer comparison, and small study groups

This target user definition fits the onboarding flow that asks users to select exam target, attempt number, and daily study hours. Those fields are not just profile inputs. They are the starting variables for personalization of study flows, mentor context, and planning intensity.

User Needs and Jobs to Be Done

The product should solve a set of concrete jobs to be done for the aspirant:

User need	Product job
Tell me what to study next	Provide a syllabus aware AI mentor and a structured notes tree
Tell me whether I am on track	Show a health check score, readiness trend, and subject progress indicators
Help me revise daily without friction	Surface a daily question and one tap quiz flow
Keep me motivated through long prep cycles	Use streaks, coins, reminders, and visible momentum
Make current affairs exam relevant	Curate current affairs with tags such as GS Paper 2 and subject filters
Reduce loneliness in self study	Offer peer pods and lightweight group discussions
Work even with weak internet	Keep the mobile interface light and include low data mode

These needs justify the feature set that emerged in the final Figma build. Each major screen maps directly to one or more of these jobs rather than existing as decorative functionality.

Product Concept

Mitra AI is best understood as a hybrid of four product models that the case study explicitly asks the PM to choose between: planning and execution system, AI driven mentor or feedback engine, community accountability product, and hybrid model. The prototype clearly evolves into a hybrid, but one with a strong center of gravity.

The center of gravity is the AI mentor plus progress loop. The mentor helps with direction, the notes system reduces resource hunting, the daily question builds habit, the progress dashboard creates feedback, current affairs make learning relevant, and pods extend motivation through community. This sequencing suggests that the product is not community first and not content first. It is mentor first, supported by execution and accountability layers.

MVP Scope

The MVP should stay disciplined and focus on features that directly solve the case study pain points. Based on the supplied prototype and the final app structure, the MVP consists of the following modules.

Onboarding

The onboarding flow has three short steps. Users select their exam target, their attempt number, and their daily study hours. A Get Started action ends the flow.

This is a smart MVP choice because it captures the minimum information needed to personalize the product without creating friction. It also frames the user relationship around guided preparation from the first interaction rather than passive browsing.

Home Dashboard

The home dashboard acts as the operating center of the app. It includes a greeting, a visible health check or readiness score, a daily question card, streak and coin indicators, current affairs shortcuts, and a quick study grid.

This screen matters because it converts a broad preparation journey into a clear daily control panel. Instead of asking the user to decide from a blank slate, the product surfaces the next useful actions.

AI Mentor Chat

The mentor chat is the core guidance surface. It includes message bubbles, a topic breadcrumb, quick reply prompts, bookmark support, and a standard chat input. The example interaction around polity and constitutional amendments shows that the mentor is meant to be syllabus aware rather than generic.

For an aspirant, this feature reduces uncertainty and lowers the effort required to get contextual help. For the product strategy, it is the clearest differentiation from static prep apps.

Notes Browser

The notes browser uses a syllabus tree structure with expandable sections such as History, Polity, Economy, Geography, and Science and Technology. Subtopics contain short previews and are searchable.

This module directly targets the collection trap from the case study. Instead of forcing users to gather PDFs across channels, the app creates a single structured repository organized around the syllabus.

Progress Dashboard

The upgraded progress area includes coverage, consistency, and readiness indicators, a fourteen day trend chart, weak topic flags, subject coverage bars, and streak visibility. This is one of the strongest responses to the feedback gap described in the case study.

Aspirants preparing for months or years need evidence that their study routine is translating into progress. Even if the early MVP metrics are approximate, visual feedback creates a much stronger daily loop than occasional test scores alone.

Current Affairs Feed

The current affairs feed includes a vertical card stream with headlines, summaries, exam relevance tags, trending labels, and subject filters. This design moves current affairs from passive consumption toward guided relevance.

That is an important product decision. Most learners do not need more news. They need exam mapped current affairs that connect current events to likely preparation value.

Daily Question and Result Flow

The daily question flow gives the user a multiple choice question, requires a selection before submission, and then shows correctness, explanation, and coins earned. This is a compact but high value loop because it combines retrieval practice, small feedback moments, and motivation.

It also works well as a repeatable habit surface. A daily question is easier to start than a full mock test and can become the app opening ritual.

Peer Pods

The pods feature adds group cards, member counts, activity indicators, and a group detail view with discussion threads. This feature is intentionally lighter than a full social network.

That is the correct MVP posture. The goal is accountability and shared momentum, not open ended community sprawl.

Settings and Accessibility

The final build includes low data mode, language preference, account details, logout, offline sync status, reminder controls, and notification sound settings. This makes the app more credible as a daily use product rather than a design exercise.

Low data mode and sync visibility are especially important because they connect directly to the stated user problem of unstable connectivity and make the product more realistic for the intended audience.

End to End Key Flows

A PM document should show not just features, but user movement between them. The important flows in Mitra AI are as follows.

Flow 1: New User Setup to First Value

1. User opens the app and completes onboarding by selecting exam target, attempt number, and daily study hours.
2. The app lands the user on the dashboard with a visible readiness score and daily question.

3. The user taps the question, attempts it, sees the result and coin reward, and begins to understand the daily loop.
4. The user can then open the mentor or notes to continue guided study.

This flow is effective because it moves quickly from setup to action. The user is not asked to browse endlessly before reaching value.

Flow 2: Direction Seeking to Topic Learning

1. User is unsure how to approach a topic.
2. User opens AI Mentor and asks a syllabus related question.
3. The mentor responds in context, supported by a visible topic breadcrumb.
4. The user can save or bookmark relevant guidance and move into notes for deeper reading.

This flow positions the mentor as the front door to study direction rather than just a support feature.

Flow 3: Awareness to Revision

1. The dashboard or progress tab flags weak areas.
2. The user sees the weak topic card and taps Revise Now.
3. The user is taken into a relevant content or revision surface.
4. Progress indicators reinforce whether revision effort is improving coverage or readiness.

This closes the feedback loop and helps the app feel like a coach rather than a passive library.

Flow 4: Current Affairs to Exam Context

1. User sees current affairs cards on the dashboard or opens the full feed.
2. User filters by subject or relevance.
3. User reads a concise summary and understands the exam mapping through tags.
4. The user can then connect the item to notes, mentor questions, or revision.

This flow is useful because it reduces the mental burden of deciding which news matters.

Flow 5: Isolation to Accountability

1. User enters the pods tab.
2. User sees active study groups and chooses one.

3. The user reads or contributes to group discussion.
4. Social visibility and shared activity improve consistency over time.

This is a measured response to the case study observation that self study often becomes isolating and demotivating.

Information Architecture

The app ends up with a multi tab structure that reflects recurring study behaviors. Home is for daily command center tasks. Mentor is for guidance. Notes is for structured content access. Progress is for self evaluation. Pods is for accountability. Settings supports personalization and device constraints.

That architecture is strong because each tab has a clear mental model. There is low overlap between surfaces, which is important for a product intended to be used daily on a phone.

Feature Rationale

Each feature in the app should exist for a reason. The rationale for the final set is below.

Feature	Why it exists
Health check score	Gives users a simple readiness signal instead of vague anxiety
Daily question	Creates a low friction habit loop and supports retrieval practice
Streaks and coins	Reinforce consistency and make effort visible
AI mentor	Reduces uncertainty and gives personalized study direction
Searchable notes tree	Cuts down resource hunting and syllabus fragmentation
Current affairs tags	Turns general news into exam relevant knowledge
Trend chart	Makes progress visible over time rather than in isolated moments
Weak topic card	Converts analytics into an immediate next action
Pods	Recreates some accountability missing in solo preparation
Low data mode	Increases accessibility for lower bandwidth users
Language settings	Makes the product more inclusive for a wider aspirant base

Differentiation vs Existing Solutions

The prototype establishes a differentiated position by combining several useful behaviors in one product rather than competing on a single learning format.

Existing market pattern	Limitation	Mitra AI response
Static notes and video libraries	Too much passive content, low personalization	AI guided mentor plus structured notes
Generic quiz apps	Weak connection to syllabus gaps and long term readiness	Daily questions plus readiness and coverage tracking
Current affairs channels	High volume, low exam context	Relevance tagged current affairs feed
Traditional heavy prep apps	Assumes strong internet and often replicates classroom models online	Lightweight mobile first design with low data support
Informal peer groups	No structure, weak moderation, scattered conversations	Purpose built pods inside the study workflow

The differentiation is strongest when described as better decision support. The app is not simply easier to use. It is intended to help the aspirant make better daily choices with less cognitive overhead.

Monetization Strategy

The monetization plan from the prototype is sensible for this category because it balances reach with depth of value.

Freemium Entry Layer

Free users get access to daily questions, limited notes, and progress tracking. This is the right acquisition model because UPSC aspirants are highly value conscious and often compare many free resources before paying.

Premium Subscription

Paid access unlocks the full AI mentor, deeper current affairs coverage, unlimited dynamic questions, and advanced peer features. This matches the areas where personalization and repeated utility are strongest, which improves the odds of sustained subscription willingness.

Gamified Rewards

Coins can be earned through daily engagement and redeemed toward premium access. This is useful because it ties monetization to behavior rather than putting a paywall in front of early habit formation.

Institutional Partnerships

Bulk licensing with coaching centers is a practical expansion path once the product proves that its mentor and progress loops improve engagement. This channel can lower customer acquisition cost and give the startup distribution credibility.

Ad Lite Option

A contextual ad supported free tier is possible, but it should remain secondary. Ads can support reach, though overuse would conflict with the product promise of clarity and focus.

Go To Market Strategy

The GTM plan is built around communities and trust rather than mass advertising. That fits both the seed stage constraint in the case study and the behavior of exam prep audiences.

Launch Audience

The first wave should focus on UPSC aspirants who are self studying, digitally active, and likely to join a beta. Tier 2 and Tier 3 city learners and working aspirants are attractive early segments because the product value proposition directly addresses their pain points.

Acquisition Channels

The supplied plan identifies Telegram, WhatsApp, Reddit, LinkedIn groups, YouTube educator partnerships, and campus or coaching outreach. These are sensible because aspirants already discover study resources and peer support through these channels.

Launch Motion

A beta launch with early adopters should come first, followed by a referral loop and content led promotion through current affairs snippets and mentor demonstrations. This sequence is stronger than a broad paid launch because the product still needs qualitative proof that users trust the mentor and return regularly.

Messaging

The strongest message is that Mitra AI is a syllabus aware, gamified, and accessible UPSC mentor. That phrasing works because it combines usefulness, motivation, and reach in one value proposition.

Risks and Trade Offs

A stronger PM document should go deeper on the risks.

AI Accuracy Risk

If the mentor gives inaccurate guidance, weak notes, or misleading explanations, user trust can collapse quickly. In this category, trust is not a nice to have feature. It is part of the product itself.

Over Gamification Risk

Coins, streaks, and rewards can drive retention, but they can also distort user behavior if they become the goal instead of disciplined learning. The system must reward meaningful preparation, not superficial taps.

Community Risk

Pods can increase consistency, but they also create moderation, misinformation, and unhealthy comparison risks. A lightweight structured community is good. An unmanaged social surface is dangerous.

Scope Risk

The prototype already contains many surfaces for an MVP. While the concept is coherent, a seed stage team would likely need to sequence development carefully to avoid spreading engineering and content operations too thin.

Accessibility Trade Off

A low data mobile first experience helps distribution, but it limits rich media possibilities such as long form video or visually heavy interactive learning. That is an acceptable trade off if the product remains focused on guidance and execution.

Adoption Barrier

Many aspirants already use known brands, coaching ecosystems, and familiar channels. The product will need to prove credibility quickly through usefulness, not just good design.

Sequencing Recommendation

Although the prototype includes many features, a PM preparing the actual launch would likely ship them in phases.

Phase 1

- Onboarding
- Home dashboard
- Daily question flow

- AI mentor
- Notes browser
- Basic progress indicators

This phase delivers the core direction and habit loop.

Phase 2

- Full progress dashboard with trends and weak topics
- Current affairs feed with relevance tags
- Low data mode and language controls

This phase strengthens retention and practical utility.

Phase 3

- Peer pods
- Expanded sync support
- More mature personalization settings

This phase adds accountability and ecosystem depth once the core individual experience is trusted.

This sequencing would reduce launch risk while preserving the final product vision shown in the Figma build.

Success Metrics

Even though the supplied materials do not define formal metrics, the product naturally suggests what should be measured.

Goal	Indicative metric
Activation	Onboarding completion rate and first daily question attempt
Habit formation	Weekly active users, streak continuation, daily return rate
Guidance value	Mentor sessions per active user and repeat mentor usage
Learning engagement	Notes opens, current affairs reads, quiz completions
Feedback trust	Progress tab revisit rate and revise now click through
Community health	Pod participation rate and low moderation incident rate
Monetization readiness	Free to premium conversion and reward redemption behavior

These metrics would help determine whether the app is becoming a true operating system or just another occasionally opened study tool.

Why the Prototype Works

The strongest part of the Mitra AI concept is that it translates abstract exam stress into manageable daily interactions. A user does not need to think in terms of a two year preparation burden every time they open the app. They can see readiness, answer one question, review one weak area, ask one mentor question, or read one relevant current affairs summary.

That compression of complexity into clear next actions is good product design. It is also exactly the kind of utility the case study asks a founding PM to create under realistic startup constraints.

Next Milestones

The next milestones from the supplied concept are directionally strong and can be made more operational.

1. Run a pilot with a small cohort of UPSC aspirants to test whether the daily loop actually drives repeat usage.
2. Validate mentor trust by reviewing answer quality, note quality, and weak topic suggestions.
3. Study whether users understand the readiness score and whether it changes behavior.
4. Observe whether pods improve consistency or simply create noise.
5. Build partnerships with coaching institutes and educators only after the product demonstrates clear value in self study workflows.
6. Expand features based on evidence from engagement, retention, and subscription intent rather than adding more surface area too early.

Final Conclusion

Mitra AI is a credible MVP concept because it starts with a sharp user problem and turns that problem into a cohesive mobile product. The combination of AI mentor guidance, structured notes, daily habit loops, progress visibility, exam relevant current affairs, and lightweight accountability maps closely to the behavioral gaps highlighted in the case study.